

HIVE CASE STUDY

Problem Statement:

With online sales gaining popularity, tech companies are exploring ways to improve their sales by analyzing customer behavior and gaining insights about product trends. Furthermore, the websites make it easier for customers to find the products they require without much scavenging. Needless to say, the role of big data analysts is among the most sought-after job profiles of this decade. Therefore, as part of this assignment, we will be challenging you, as a big data analyst, to extract data and gather insights from a real-life data set of an e-commerce company.

The implementation phase can be divided into the following parts:

- Copying the data set into the HDFS:
- Launch an EMR cluster that utilizes the Hive services, and
- Move the data from the S3 bucket into the HDFS
- Creating the database and launching Hive queries on your EMR cluster:
- Create the structure of your database,
- Use optimized techniques to run your queries as efficiently as possible
- Show the improvement of the performance after using optimization on any single query.
- Run Hive queries to answer the questions given below.
- Cleaning up -:
- Drop your database, and
- Terminate your cluster

The data is available from the link provided:

<https://e-commerce-events-ml.s3.amazonaws.com/2019-Oct.csv>

<https://e-commerce-events-ml.s3.amazonaws.com/2019-Nov.csv>

Overview of steps:

- **Copying the data set into the HDFS:**

- Launch an EMR cluster that utilizes the Hive services, and
- Move the data from the S3 bucket into the HDFS

- **Creating the database and launching Hive queries on your EMR cluster:**

- Create the structure of your database,
- Use optimized techniques to run your queries as efficiently as possible
- Show the improvement of the performance after using optimization on any single query.
- Run Hive queries to answer the questions given below.

- **Cleaning up**

- Drop your database, and
- Terminate your cluster

Data Collection and Processing:

EMR Cluster Creation

Create Cluster - Advanced Options [Go to quick options](#)

Step 1: Software and Steps
Step 2: Hardware
Step 3: General Cluster Settings
Step 4: Security

Software Configuration

Release emr-5.36.0

☒ Hadoop 2.10.1
☐ JupyterHub 1.4.1
☐ Ganglia 3.7.2
☒ Hive 2.3.9
☐ JupyterEnterpriseGateway 2.1.0
☐ Mahout 0.13.0
☐ Oozie 5.2.1
☐ TensorFlow 2.4.1

☐ Zeppelin 0.10.0
☐ Tez 0.9.2
☐ HBase 1.4.13
☐ Presto 0.267
☐ MXNet 1.8.0
☒ Hue 4.10.0
☒ Spark 2.4.8

☐ Livy 0.7.1
☐ Flink 1.14.2
☒ Pig 0.17.0
☐ ZooKeeper 3.4.14
☐ Sqoop 1.4.7
☐ Phoenix 4.14.3
☐ HCatalog 2.3.9

Hardware Configuration Page > To define the cluster & nodes : Instance type for both master & core nodes are M4.large

Node type	Instance type	Instance count	Purchasing option
Master Master - 1	m4.large 2 vCore, 8 GiB memory, EBS only storage EBS Storage: 32 GiB Add configuration settings	1 Instances	<input checked="" type="radio"/> On-demand <input type="radio"/> Spot Use on-demand as max price
Core Core - 2	m4.large 2 vCore, 8 GiB memory, EBS only storage EBS Storage: 32 GiB Add configuration settings	2 Instances	<input checked="" type="radio"/> On-demand <input type="radio"/> Spot Use on-demand as max price
Task Task - 3	m5.xlarge 4 vCore, 16 GiB memory, EBS only storage EBS Storage: 32 GiB Add configuration settings	0 Instances	<input checked="" type="radio"/> On-demand <input type="radio"/> Spot Use on-demand as max price

[+ Add task instance group](#)

Amazon EMR

EMR Studio

EMR Serverless New

EMR on EC2

Clusters

Notebooks

Git repositories

Security configurations

Block public access

VPC subnets

Events

EMR on EKS

Virtual clusters

Help

What's new

EMR Serverless is now GA.

With EMR Serverless, get the benefits of Amazon EMR such as open source compatibility, latest versions and performance optimized runtime for popular frameworks along with easy provisioning, quick job startup, automatic capacity management, and simple cost controls. [Get Started with EMR Serverless.](#)

Clone Terminate AWS CLI export

Cluster: Hive case study Starting Configuring cluster software

Summary Application user interfaces Monitoring Hardware Configurations Events Steps Bootstrap actions

Summary

ID: j-1OFFCUHILOVFZ

Creation date: 2022-08-02 12:38 (UTC+5:30)

Elapsed time: 3 minutes

After last step completes: Cluster waits

Termination protection: On [Change](#)

Tags: -- [View All / Edit](#)

Master public DNS: ec2-44-205-19-17.compute-1.amazonaws.com [Connect to the Master Node Using SSH](#)

Configuration details

Release label: emr-5.36.0

Hadoop distribution: Amazon 2.10.1

Applications: Hive 2.3.9, Pig 0.17.0, Hue 4.10.0, Spark 2.4.8

Log URI: s3://aws-logs-751171450291-us-east-1/elasticmapreduce/

EMRFS consistent view: Disabled

Custom AMI ID: --

Amazon Linux Release: 2.0.20220426.0 [Learn more](#)

Application user interfaces

Persistent user interfaces: --

On-cluster user Not Enabled [Enable an SSH Connection](#)

Interfaces:

Network and hardware

Availability zone: us-east-1d

Subnet ID: [subnet-08f6ad6e](#)

Master: Bootstrapping 1 m4.large

Core: Provisioning 2 m4.large

Task: --

Cluster scaling: Not enabled

Auto-termination: Terminate if idle for 1 hour

Hadoop & Hive Queries :

Terminal > Connecting to EMR Cluster using ssh.

```

Last login: Tue Aug  2 22:23:37 on ttys000

The default interactive shell is now zsh.
To update your account to use zsh, please run `chsh -s /bin/zsh`.
For more details, please visit https://support.apple.com/kb/HT208050.
(base) roshans-MacBook-Air:~ roshanc$ ssh -i ~/keypair.pem -ND 8157 hadoop@ec2-3-231-60-23.compute-1.amazonaws.com
ssh: connect to host ec2-3-231-60-23.compute-1.amazonaws.com port 22: Operation timed out
(base) roshans-MacBook-Air:~ roshanc$ ssh -i ~/keypair.pem -ND 8157 hadoop@ec2-3-231-60-23.compute-1.amazonaws.com
ssh: Could not resolve hostname ec2-3-231-60-23.compute-1.amazonaws.com: nodename nor servname provided, or not known
(base) roshans-MacBook-Air:~ roshanc$ ssh -i ~/keypair.pem hadoop@ec2-3-231-60-23.compute-1.amazonaws.com
ssh: connect to host ec2-3-231-60-23.compute-1.amazonaws.com port 22: Operation timed out
(base) roshans-MacBook-Air:~ roshanc$ ssh -i keypair.pem hadoop@ec2-3-231-60-23.compute-1.amazonaws.com
^C
(base) roshans-MacBook-Air:~ roshanc$
(base) roshans-MacBook-Air:~ roshanc$ pwd
/Users/roshanc
(base) roshans-MacBook-Air:~ roshanc$ ssh -i ~/keypair.pem hadoop@ec2-3-231-60-23.compute-1.amazonaws.com
ssh: connect to host ec2-3-231-60-23.compute-1.amazonaws.com port 22: Operation timed out
(base) roshans-MacBook-Air:~ roshanc$ ssh -i ~/keypair.pem hadoop@ec2-3-231-60-23.compute-1.amazonaws.com
The authenticity of host 'ec2-3-231-60-23.compute-1.amazonaws.com (3.231.60.23)' can't be established.
ECDSA key fingerprint is SHA256:tcfN8N1T09BojRpnXpmZ/6qUAdrAhSkWbGk0Tc0FMWE.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-3-231-60-23.compute-1.amazonaws.com,3.231.60.23' (ECDSA) to the list of known hosts.
Last login: Wed Aug  3 08:19:48 2022

 _ | _ | _ )
 _ | ( _ | /
 _ | \ _ | _ |

Amazon Linux 2 AMI

https://aws.amazon.com/amazon-linux-2/
-bash: warning: setlocale: LC_CTYPE: cannot change locale (UTF-8): No such file or directory

EEEEEEEEEEEEEEEEEEEE MMMMMMM      MMMMMMM RRRRRRRRRRRRRRR
E::::::::::::::::::::E M:::::::::M      M:::::::::M R:::::::::::::R
EE::::::::EEEEEEEE::::E M:::::::::M      M:::::::::M R::::::::RRRRR:::::R
  E::::E      EEEEE  M:::::::::M      M:::::::::M RRR::::R      R::::R
  E::::E      M:::::M M::::M::::M      R:::R      R::::R
  E::::EEEEEEEEEE  M:::::M M::::M::::M      R::RRRRR:::::R
  E:::::::::::::E  M:::::M M::::M::::M      R:::::::::RR
  E::::EEEEEEEEEE  M:::::M M:::::M      R::RRRRR:::::R
  E::::E      M:::::M M::::M      R:::R      R::::R
  E::::E      EEEEE  M:::::M      M:::::M      R::::R      R::::R
EE::::::::EEEEEEEE::::E M:::::M      M:::::M      R::::R      R::::R
E:::::::::::::E M:::::M      M:::::M RRR::::R      R::::R
EEEEEEEEEEEEEEEEEEEE MMMMMMM      MMMMMMM RRRRRRR      RRRRRR

[hadoop@ip-172-31-4-81 ~]$

```

Creating a directory "casestudy"

```
hadoop fs -mkdir /
```

```
casestudyhadoop fs -ls /
```

```
Found 5 items
drwxr-xr-x - hdfs hdfsadmingroup 0 2022-08-03 08:08 /apps
drwxr-xr-x - hadoop hdfsadmingroup 0 2022-08-03 08:29 /casestudy
drwxrwxrwt - hdfs hdfsadmingroup 0 2022-08-03 08:10 /tmp
drwxr-xr-x - hdfs hdfsadmingroup 0 2022-08-03 08:08 /user
drwxr-xr-x - hdfs hdfsadmingroup 0 2022-08-03 08:08 /var
[hadoop@ip-172-31-4-81 ~]$
```

Loading the datasets into HDFS from S3:

```
hadoop distcp 's3://e-commerce-events-ml/2019-Oct.csv' /casestudy/2019_Oct.csv
```

```
File Input Format Counters
  Bytes Read=238
File Output Format Counters
  Bytes Written=0
DistCp Counters
  Bytes Copied=482542278
  Bytes Expected=482542278
  Files Copied=1
```

```
hadoop distcp 's3://e-commerce-events-ml/2019-Nov.csv' /casestudy/2019_Nov.csv
```

```
File Input Format Counters
  Bytes Read=238
File Output Format Counters
  Bytes Written=0
DistCp Counters
  Bytes Copied=545839412
  Bytes Expected=545839412
  Files Copied=1
```

Viewing the data

```
hadoop fs -cat /casestudy/2019_Oct.csv |head
```

```
[hadoop@ip-172-31-4-81 ~]$ hadoop fs -cat /casestudy/2019_Oct.csv |head
event_time,event_type,product_id,category_id,category_code,brand,price,user_id,user_session
2019-10-01 00:00:00 UTC,cart,5773203,1487580005134238553,,runail,2.62,463240011,26dd6e6e-4dac-4778-8d2c-92e149dab885
2019-10-01 00:00:03 UTC,cart,5773353,1487580005134238553,,runail,2.62,463240011,26dd6e6e-4dac-4778-8d2c-92e149dab885
2019-10-01 00:00:07 UTC,cart,5881589,2151191071051219817,,lovely,13.48,429681830,49e8d843-adf3-428b-a2c3-fe8bc6a307c9
2019-10-01 00:00:07 UTC,cart,5723490,1487580005134238553,,runail,2.62,463240011,26dd6e6e-4dac-4778-8d2c-92e149dab885
2019-10-01 00:00:15 UTC,cart,5881449,1487580013522845895,,lovely,0.56,429681830,49e8d843-adf3-428b-a2c3-fe8bc6a307c9
2019-10-01 00:00:16 UTC,cart,5857269,1487580005134238553,,runail,2.62,430174032,73dea1e7-664e-43f4-8b30-d32b9d5af04f
2019-10-01 00:00:19 UTC,cart,5739055,1487580008246412266,,kapous,4.75,377667011,81326ac6-daa4-4f0a-b488-fd0956a78733
2019-10-01 00:00:24 UTC,cart,5825598,1487580009445982239,,0.56,467916806,2f5b5546-b8cb-9ee7-7ecd-84276f8ef486
2019-10-01 00:00:25 UTC,cart,5698989,1487580006317032337,,1.27,385985999,d30965e8-1101-44ab-b45d-cc1bb9fae694
```

```
hadoop fs -cat /casestudy/2019_Nov.csv | head
```

```
event_time,event_type,product_id,category_id,category_code,brand,price,user_id,user_session
2019-11-01 00:00:02 UTC,view,5802432,1487580009286598681,,0.32,562076640,09fafd6c-6c99-46b1-834f-33527f4de241
2019-11-01 00:00:09 UTC,cart,5844397,1487580006317032337,,2.38,553329724,2067216c-31b5-455d-a1cc-af0575a34ffb
2019-11-01 00:00:10 UTC,view,5837166,1783999064103190764,,pnb,22.22,556138645,57ed222e-a54a-4907-9944-5a875c2d7f4f
2019-11-01 00:00:11 UTC,cart,5876812,1487580010100293687,,jessnail,3.16,564506666,186c1951-8052-4b37-adce-dd9644b1d5f7
2019-11-01 00:00:24 UTC,remove_from_cart,5826182,1487580007483048900,,3.33,553329724,2067216c-31b5-455d-a1cc-af0575a34ffb
2019-11-01 00:00:24 UTC,remove_from_cart,5826182,1487580007483048900,,3.33,553329724,2067216c-31b5-455d-a1cc-af0575a34ffb
2019-11-01 00:00:25 UTC,view,5856189,1487580009026551821,,runail,15.71,562076640,09fafd6c-6c99-46b1-834f-33527f4de241
2019-11-01 00:00:32 UTC,view,5837835,1933472286753424063,,3.49,514649199,432a4e95-375c-4b40-bd36-0fc039e77580
2019-11-01 00:00:34 UTC,remove_from_cart,5870838,1487580007675986893,,milv,0.79,429913900,2f0bfff3c-252f-4fe6-afcd-5d8a6a92839a
```

Datasets are successfully loaded.

Launch Hive

```
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.properties Async: false
hive> show databases ;
OK
default
Time taken: 1.895 seconds, Fetched: 1 row(s)
```

Creating new database “hive_assignment”

```
hive> CREATE DATABASE IF NOT EXISTS hive_assignment ;
hive> SHOW DATABASES ;
hive> DESCRIBE DATABASE hive_assignment ;
```

```
hive> CREATE DATABASE IF NOT EXISTS hive_assignment ;
OK
Time taken: 0.077 seconds
hive> SHOW DATABASES ;
OK
default
hive_assignment
Time taken: 0.031 seconds, Fetched: 2 row(s)
hive> DESCRIBE DATABASE hive_assignment ;
OK
hive_assignment      hdfs://ip-172-31-4-81.ec2.internal:8020/user/hive/warehouse/hive_assignment.db  hadoop  USER
Time taken: 0.048 seconds, Fetched: 1 row(s)
```

Creating new table “retail”

```
hive > CREATE EXTERNAL TABLE IF NOT EXISTS retail (event_time timestamp, event_type string,
product_id string, category_id string, category_code string, brand string, price decimal(10,3), user_id
bigint, user_session string) ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde'
WITH SERDEPROPERTIES ("separatorChar" = ",", "quoteChar" = "\"", "escapeChar" = "\\") stored as
textfile LOCATION '/casestudy' TBLPROPERTIES ("skip.header.line.count"="1") ;
```

```
hive> CREATE EXTERNAL TABLE IF NOT EXISTS retail (event_time timestamp, event_type string, product_id string, category_id string,
category_code string, brand string, price decimal(10,3), user_id bigint, user_session string) ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde' WITH SERDEPROPERTIES ("separatorChar" = ",", "quoteChar" = "\"", "escapeChar" = "\\") stored as
textfile LOCATION '/casestudy' TBLPROPERTIES ("skip.header.line.count"="1") ;
OK
Time taken: 0.395 seconds
```

```
hive> DESCRIBE retail ;
```

```
hive> DESCRIBE retail ;
OK
event_time          string              from deserializer
event_type          string              from deserializer
product_id          string              from deserializer
category_id         string              from deserializer
category_code       string              from deserializer
brand               string              from deserializer
price               string              from deserializer
user_id             string              from deserializer
user_session        string              from deserializer
Time taken: 0.091 seconds, Fetched: 9 row(s)
hive> █
```

Loading data into table “retail”;

```
hive> LOAD DATA INPATH '/casestudy/2019_Oct.csv' INTO TABLE retail ;
```

```
hive> LOAD DATA INPATH '/casestudy/2019_Nov.csv' INTO TABLE retail;
```

```
[hive> LOAD DATA INPATH '/casestudy/2019_Oct.csv' INTO TABLE retail ;
Loading data to table default.retail
OK
Time taken: 0.482 seconds
[hive> LOAD DATA INPATH '/casestudy/2019_Nov.csv' INTO TABLE retail;
Loading data to table default.retail
OK
Time taken: 0.521 seconds
hive> ]
```

Performing data check:

```
hive> SELECT * FROM retail WHERE MONTH(event_time)=11 limit 5;
```

```
hive> SELECT * FROM retail WHERE MONTH(event_time)=10 limit 5 ;
```

```
2019-10-01 00:00:00 UTC cart 5773203 1487580005134238553 runail 2.62 463240011 26dd6e6e-4dac-4778-8d2c-
92e149dab885
2019-10-01 00:00:03 UTC cart 5773353 1487580005134238553 runail 2.62 463240011 26dd6e6e-4dac-4778-8d2c-
92e149dab885
2019-10-01 00:00:07 UTC cart 5881589 2151191071051219817 lovely 13.48 429681830 49e8d843-adf3-428b-a2c3-
fe8bc6a307c9
2019-10-01 00:00:07 UTC cart 5723490 1487580005134238553 runail 2.62 463240011 26dd6e6e-4dac-4778-8d2c-
92e149dab885
2019-10-01 00:00:15 UTC cart 5881449 1487580013522845895 lovely 0.56 429681830 49e8d843-adf3-428b-a2c3-
fe8bc6a307c9
Time taken: 0.232 seconds, Fetched: 5 row(s)
```

You are required to provide answers to the questions given below

Q) Find the total revenue generated due to purchases made in October.

```
hive> SELECT SUM(price) FROM retail WHERE MONTH(event_time)=10 AND event_type='purchase' ;
```

```
[hive> SELECT SUM(price) FROM retail WHERE MONTH(event_time)=10 AND event_type='purchase' ;
Query ID = hadoop_20220803085430_095fefda-88a2-4476-8899-59518822d3bd
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1659514121071_0004)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	2	2	0	0	0	0	
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0	

VERTICES: 02/02 [=====] 100% ELAPSED TIME: 40.85 s

```
OK
1211538.4299997438
Time taken: 50.638 seconds, Fetched: 1 row(s)
hive> set hive.exec.dynamic.partition=true;
hive> set hive.exec.dynamic.partition.mode=nonstrict;
hive> ]
```

Time taken to execute the above query is 40.85 sec.

DYNAMIC PARTITIONING

```
hive> set hive.exec.dynamic.partition=true;
```

```
hive> set hive.exec.dynamic.partition.mode=nonstrict;
```


PARTITION TABLE 1: retail_part_1

Partition on : event_type (there are 4 types and all questions are related to 'purchase')

```
hive> CREATE EXTERNAL TABLE IF NOT EXISTS retail_part_1 (event_time timestamp, product_id string,
category_id string, category_code string, brand string, price decimal(10,3), user_id bigint, user_session
string) PARTITIONED BY(event_type string) CLUSTERED BY (user_id) INTO 5 buckets ROW FORMAT
SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde' STORED AS textfile ;
```

```
hive> DESCRIBE retail_part_1 ;
```

```
hive> DESCRIBE retail_part_1 ;
OK
event_time          string              from deserializer
product_id          string              from deserializer
category_id         string              from deserializer
category_code       string              from deserializer
brand               string              from deserializer
price               string              from deserializer
user_id             string              from deserializer
user_session        string              from deserializer
event_type          string
# Partition Information
# col_name          data_type          comment
event_type          string
Time taken: 0.106 seconds, Fetched: 14 row(s)
hive>
```

```
hive> INSERT INTO TABLE retail_part_1 PARTITION (event_type) SELECT event_time, product_id,
category_id, category_code, brand, price, user_id, user_session, event_type FROM retail ;
```

```
hive> SELECT SUM(price) FROM retail_part_3 WHERE MONTH(event_time)=10 AND event_type='purchase' ;
Query ID = hadoop_20220803091803_ac5f75c4-2cce-47ee-9750-5e3abd557fe9
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1659514121071_0007)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1		container	SUCCEEDED	0	0	0	0	0	0
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 01/02 [=====>>>] 100% ELAPSED TIME: 5.91 s
OK
NULL
Time taken: 14.353 seconds, Fetched: 1 row(s)
hive>
```

Executing the same query with the new table "retail_part_1" to check the time.

```
hive> SELECT SUM(price) FROM retail_part_1 WHERE MONTH(event_time)=10
AND event_type='purchase' ;
```

```

Query ID = hadoop_20220803090114_551991c9-cc8e-4c53-8500-28ecb30a9f42
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1659514121071_0005)

```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container		SUCCEEDED	2	2	0	0	0	0
Reducer 2	container		SUCCEEDED	5	5	0	0	0	0

```

VERTICES: 02/02 [=====] 100% ELAPSED TIME: 133.27 s

```

```

Loading data to table default.retail_part_1 partition (event_type=null)

```

```

Loaded : 4/4 partitions.
Time taken to load dynamic partitions: 0.377 seconds
Time taken for adding to write entity : 0.001 seconds

```

```

OK
Time taken: 142.523 seconds
hive>

```

PARTITION TABLE 2: retail_part_3

Partition on : month

```

hive> CREATE EXTERNAL TABLE IF NOT EXISTS retail_part_3 (event_time timestamp, event_type string,
product_id string, category_id string, category_code string, brand string, price decimal(10,3), user_id bigint,
user_session string) PARTITIONED BY(month int) CLUSTERED BY (brand) INTO 5 buckets ROW FORMAT
SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde' STORED AS textfile ;

```

```

hive> DESCRIBE retail_part_3 ;

```

```

hive> DESCRIBE retail_part_3 ;
OK
event_time          string              from deserializer
event_type          string              from deserializer
product_id          string              from deserializer
category_id         string              from deserializer
category_code       string              from deserializer
brand               string              from deserializer
price               string              from deserializer
user_id             string              from deserializer
user_session        string              from deserializer
month               int
# Partition Information
# col_name          data_type          comment
month               int
Time taken: 0.057 seconds, Fetched: 15 row(s)
hive>

```

Executing the same query with the new table “retail_part_3” to check the time.

```

hive> SELECT SUM(price) FROM retail_part_3 WHERE MONTH(event_time)=10
AND event_type='purchase' ;

```



```

Time taken: 0.1316 seconds
hive> SELECT SUM(price) FROM retail_part_1 WHERE MONTH(event_time)=10 AND event_type='purchase' ;
Query ID = hadoop_20220803090901_a5596916-93ed-4b10-bf17-ba2790192d2f
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1659514121071_0006)

```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	5	5	0	0	0	0	
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0	

```

VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 24.20 s
OK
1211538.4300000467
Time taken: 35.347 seconds, Fetched: 1 row(s)

```

We get an optimised table by Partitioning on “event_type” and clustering by “user_id” . Hence, for all the following analysis, we will be using the optimised table “retail_part_1”.

Q) Find the total revenue generated due to purchases made in October.

```
hive> SELECT SUM(price) FROM retail_part_1 WHERE MONTH(event_time)=10  
AND event_type='purchase' ;
```

```
hive> SELECT SUM(price) FROM retail_part_1 WHERE MONTH(event_time)=10 AND event_type='purchase' ;  
Query ID = hadoop_20220803091950_42d4c487-9016-4223-b202-6fab445e31b2  
Total jobs = 1  
Launching Job 1 out of 1  
Status: Running (Executing on YARN cluster with App id application_1659514121071_0007)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	5	5	0	0	0	0	
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0	

```
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 23.61 s  
OK  
1211538.4300000467  
Time taken: 24.556 seconds, Fetched: 1 row(s)  
hive>
```

Time Taken to execute the above query is 25.36 sec.

Q) Write a query to yield the total sum of purchases per month in a single output.

```
hive> SELECT MONTH(event_time), SUM(price) as sum_purchase, COUNT(event_type) as cnt  
FROM retail_part_1 WHERE event_type='purchase' GROUP BY MONTH(event_time) ;
```

```
hive> SELECT MONTH(event_time), SUM(price) as sum_purchase, COUNT(event_type) as cnt FROM retail_part_1 WHERE event_type='purchase' GROUP BY MONTH(event_time) ;  
Query ID = hadoop_20220803092106_7ecdefa4-bcf1-487a-a899-62a1a1d82f86  
Total jobs = 1  
Launching Job 1 out of 1  
Status: Running (Executing on YARN cluster with App id application_1659514121071_0007)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	5	5	0	0	0	0	
Reducer 2	container	SUCCEEDED	2	2	0	0	0	0	

```
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 23.47 s  
OK  
10 1211538.4300000465 245624  
11 1531016.8999999743 322417  
Time taken: 24.156 seconds, Fetched: 2 row(s)  
hive>
```

In October month, 245624 purchases generated revenue of 1211538. Similarly in November month, 322417 purchases generated revenue of 1531016.8

Q) Write a query to find the change in revenue generated due to purchases from October to November.

```
hive>WITH diff AS ( SELECT SUM(CASE WHEN date_format(event_time,'MM')=10 THEN price ELSE 0 END) AS October, SUM(CASE WHEN date_format(event_time,'MM')=11 THEN price ELSE 0 END) AS November FROM retail_part_1 WHERE date_format(event_time,'MM') IN (10,11) AND event_type='purchase') SELECT October, November, (November - October) as Difference FROM diff ;
```

```
hive> WITH diff AS ( SELECT SUM(CASE WHEN date_format(event_time,'MM')=10 THEN price ELSE 0 END) AS October, SUM(CASE WHEN date_format(event_time,'MM')=11 THEN price ELSE 0 END) AS November FROM retail_part_1 WHERE date_format(event_time,'MM') IN (10,11) AND event_type='purchase') SELECT October, November, (November - October) as Difference FROM diff ;
Query ID = hadoop_20220803092222_1014abed-038f-410c-b54b-9024eb876abc
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1659514121071_0007)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	5	5	0	0	0	0
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 02/02 [=====] 100% ELAPSED TIME: 33.95 s
OK
1211538.4300000467      1531016.8999999743      319478.46999992756
Time taken: 34.541 seconds, Fetched: 1 row(s)
hive>
```

The change in revenue generated from October to November is 319478

Q) Find distinct categories of products. Categories with null category code can be ignored.

```
hive>SELECT DISTINCT split(category_code,'\\.[0] AS category FROM retail_part_1 WHERE split(category_code,'\\.[0]<>' ;
```

```
hive> SELECT DISTINCT split(category_code,'\\.[0] AS category FROM retail_part_1 WHERE split(category_code,'\\.[0]<>' ;
Query ID = hadoop_20220803092432_5f65e3f6-e579-44d1-8015-f1ea5b4504a6
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1659514121071_0007)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	14	14	0	0	0	0
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 02/02 [=====] 100% ELAPSED TIME: 55.58 s
OK
accessories
apparel
appliances
furniture
sport
stationery
Time taken: 56.54 seconds, Fetched: 6 row(s)
hive>
```

There are 6 distinct categories of products. They are: Furniture, appliances, accessories, apparel, sport and stationary.

Q) Find the total number of products available under each category.

```
hive>SELECT split(category_code,'\\.')[0] AS category, COUNT(product_id) AS prd FROM retail_part_1  
GROUP BY split(category_code,'\\.')[0] ORDER BY prd DESC ;
```

```
hive> SELECT split(category_code,'\\.')[0] AS category, COUNT(product_id) AS prd FROM retail_part_1 GROUP BY split(category_code, '\\.')[0] ORDER BY prd DESC ;
Query ID = hadoop_20220803092721_eb63d19b-c458-4711-b41b-fdff1c738992
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1659514121071_0007)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	14	14	0	0	0	0	0
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0	0
Reducer 3	container	SUCCEEDED	1	1	0	0	0	0	0

```
VERTICES: 03/03 [=====] 100% ELAPSED TIME: 56.11 s
OK
      8594895
appliances      61736
stationery      26722
furniture       23604
apparel 18232
accessories     12929
sport           2
Time taken: 56.727 seconds, Fetched: 7 row(s)
hive>
```

The sport category has the least number of products, whereas appliances has 61736 products.

Q) Which brand had the maximum sales in October and November combined?

```
SELECT brand, SUM(price) AS Sales FROM retail_part_1 WHERE brand <>' ' AND  
event_type='purchase'GROUP BY brand ORDER BY Sales DESC LIMIT 1 ;
```

```
hive>
> SELECT brand, SUM(price) AS Sales FROM retail_part_1 WHERE brand <>' ' AND event_type='purchase' GROUP BY brand ORDER BY Sales DESC LIMIT 1 ;
Query ID = hadoop_20220803092933_d21d1990-9269-444a-a2c8-cae0eee4157c
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1659514121071_0007)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	5	5	0	0	0	0	0
Reducer 2	container	SUCCEEDED	2	2	0	0	0	0	0
Reducer 3	container	SUCCEEDED	1	1	0	0	0	0	0

```
VERTICES: 03/03 [=====] 100% ELAPSED TIME: 21.90 s
OK
runail 148297.93999999858
Time taken: 22.512 seconds, Fetched: 1 row(s)
hive>
```

Runail has the maximum sales for both months combined.

Q)Your company wants to reward the top 10 users of its website with a Golden Customer plan.Write a query to generate a list of top 10 users who spend the most.

hive>SELECT user_id, SUM(price) AS expense FROM retail_part_1 WHERE event_type='purchase'
GROUP BY user_id ORDER BY expense DESC LIMIT 10 ;

```
hive> SELECT user_id, SUM(price) AS expense FROM retail_part_1 WHERE event_type='purchase' GROUP BY user_id ORDER BY expense DESC LIMIT 10 ;
Query ID = hadoop_20220803093522_6f677a46-c0e0-433c-a7f3-e2aee9d85b56
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1659514121071_0008)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	5	5	0	0	0	0
Reducer 2	container	SUCCEEDED	2	2	0	0	0	0
Reducer 3	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 03/03 [=====] 100% ELAPSED TIME: 28.83 s
OK
557790271      2715.8699999999991
150318419      1645.9700000000005
562167663      1352.8499999999995
531900924      1329.4499999999996
557850743      1295.4799999999996
522130011      1185.3900000000003
561592095      1109.7000000000007
431950134      1097.5899999999997
566576008      1056.3600000000006
521347209      1040.9099999999999
Time taken: 37.63 seconds, Fetched: 10 row(s)
```

Above we can find the list of the top 10 users who have spend the most.

Cleaning Up:

Once the analysis is completed, deleting the database & terminating the cluste

The screenshot shows the AWS Management Console for an Amazon EMR cluster. The cluster is named 'Hive_case_study' and is in a 'Terminated' state. The console displays various configuration details including creation date, release label, Hadoop distribution, applications, log URI, EMRFS consistent view, custom AMI ID, Amazon Linux Release, application user interfaces, persistent user interfaces, network and hardware settings, and security and access configurations.

Cluster: Hive_case_study (Terminated)

Summary:

- ID: j-11MXNGR7UTD00
- Creation date: 2022-08-03 13:30 (UTC+5:30)
- End date: 2022-08-03 15:15 (UTC+5:30)
- Elapsed time: 1 hour, 45 minutes
- After last step completes: Cluster waits
- Termination protection: Off
- Tags: --
- Master public DNS: ec2-3-231-60-23.compute-1.amazonaws.com

Configuration details:

- Release label: emr-5.36.0
- Hadoop distribution: Amazon 2.10.1
- Applications: Hive 2.3.9, Pig 0.17.0, Hue 4.10.0, Spark 2.4.8
- Log URI: s3://aws-logs-751171450291-us-east-1/elasticmapreduce/
- EMRFS consistent view: Disabled
- Custom AMI ID: --
- Amazon Linux Release: 2.0.20220426.0

Application user interfaces:

- Persistent user interfaces: Spark history server, YARN timeline server, Tez UI
- On-cluster user interfaces: --

Network and hardware:

- Availability zone: us-east-1d
- Subnet ID: subnet-08fbad6e
- Master: Terminated 1 m4.large
- Core: Terminated 2 m4.large
- Task: --
- Cluster scaling: Not enabled
- Auto-termination: Terminate if idle for 1 hour

Security and access:

- Key name: keypair
- EC2 instance profile: EMR_EC2_DefaultRole
- EMR role: EMR_DefaultRole
- Auto Scaling role: EMR_AutoScaling_DefaultRole
- Visible to all users: All
- Security groups for Master: sg-04c36ab0a89b31e0d (ElasticMapReduce-master)
- Security groups for Core & Task: sg-0ec906b09d038e1d5 (ElasticMapReduce-slave)